

✓ Predict stock prices with Long short-term memory (LSTM)

This simple example will show you how LSTM models predict time series data. Stock market data is a great choice for this because it's quite regular and widely available via the Internet.

✓ Install requirements

We install Tensorflow 2.0 with GPU support first

```
!pip install tensorflow-gpu==2.0.0-alpha0
```

```
Requirement already satisfied: tensorflow-gpu==2.0.0-alpha0 in /usr/local/lib/python3.6/dist-packages (2.0.0a0)
Requirement already satisfied: absl-py>=0.7.0 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0) (0.7.1)
Requirement already satisfied: astor>=0.6.0 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0) (0.7.1)
Requirement already satisfied: google-pasta>=0.1.2 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0) (0.1.5)
Requirement already satisfied: gast>=0.2.0 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0) (0.2.2)
Requirement already satisfied: keras-preprocessing>=1.0.5 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0)
Requirement already satisfied: tb-nightly<1.14.0a20190302,>=1.14.0a20190301 in /usr/local/lib/python3.6/dist-packages (from tensorflow-g
Requirement already satisfied: tf-estimator-nightly<1.14.0.dev2019030116,>=1.14.0.dev2019030115 in /usr/local/lib/python3.6/dist-package
Requirement already satisfied: wheel>=0.26 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0) (0.33.1)
Requirement already satisfied: keras-applications>=1.0.6 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0) (
Requirement already satisfied: protobuf>=3.6.1 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0) (3.7.1)
Requirement already satisfied: six>=1.10.0 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0) (1.11.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0) (1.1.0)
Requirement already satisfied: grpcio>=1.8.6 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0) (1.15.0)
Requirement already satisfied: numpy<2.0,>=1.14.5 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==2.0.0-alpha0) (1.16.2)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.6/dist-packages (from tb-nightly<1.14.0a20190302,>=1.14.0a20190
Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.6/dist-packages (from tb-nightly<1.14.0a20190302,>=1.14.0a201
Requirement already satisfied: h5py in /usr/local/lib/python3.6/dist-packages (from keras-applications>=1.0.6->tensorflow-gpu==2.0.0-alp
Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (from protobuf>=3.6.1->tensorflow-gpu==2.0.0-alpha0)
```

```
!pip install pandas-datareader
```

```
Requirement already satisfied: pandas-datareader in /usr/local/lib/python3.6/dist-packages (0.7.0)
Requirement already satisfied: pandas>=0.19.2 in /usr/local/lib/python3.6/dist-packages (from pandas-datareader) (0.23.4)
Requirement already satisfied: lxml in /usr/local/lib/python3.6/dist-packages (from pandas-datareader) (4.2.6)
Requirement already satisfied: wrapt in /usr/local/lib/python3.6/dist-packages (from pandas-datareader) (1.10.11)
Requirement already satisfied: requests>=2.3.0 in /usr/local/lib/python3.6/dist-packages (from pandas-datareader) (2.18.4)
Requirement already satisfied: numpy>=1.9.0 in /usr/local/lib/python3.6/dist-packages (from pandas>=0.19.2->pandas-datareader) (1.16.2)
Requirement already satisfied: python-dateutil>=2.5.0 in /usr/local/lib/python3.6/dist-packages (from pandas>=0.19.2->pandas-datareader)
Requirement already satisfied: pytz>=2011k in /usr/local/lib/python3.6/dist-packages (from pandas>=0.19.2->pandas-datareader) (2018.9)
Requirement already satisfied: urllib3<1.23,>=1.21.1 in /usr/local/lib/python3.6/dist-packages (from requests>=2.3.0->pandas-datareader)
Requirement already satisfied: idna<2.7,>=2.5 in /usr/local/lib/python3.6/dist-packages (from requests>=2.3.0->pandas-datareader) (2.6)
Requirement already satisfied: chardet<3.1.0,>=3.0.2 in /usr/local/lib/python3.6/dist-packages (from requests>=2.3.0->pandas-datareader)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.6/dist-packages (from requests>=2.3.0->pandas-datareader) (2
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.6/dist-packages (from python-dateutil>=2.5.0->pandas>=0.19.2->pandas-d
```

```
!apt install graphviz
```

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
graphviz is already the newest version (2.40.1-2).
The following package was automatically installed and is no longer required:
  libnvidia-common-410
Use 'apt autoremove' to remove it.
0 upgraded, 0 newly installed, 0 to remove and 6 not upgraded.
```

```
!pip install pydot pydot-ng
```

```
Requirement already satisfied: pydot in /usr/local/lib/python3.6/dist-packages (1.3.0)
Requirement already satisfied: pydot-ng in /usr/local/lib/python3.6/dist-packages (2.0.0)
Requirement already satisfied: pyparsing>=2.1.4 in /usr/local/lib/python3.6/dist-packages (from pydot) (2.4.0)
```

✓ Introduction

LSTMs are very powerful in sequence prediction problems. They can store past information.

▼ Loading the dataset

I use pandas-datareader to get the historical stock prices from Yahoo! finance. For this example, I get only the historical data till the end of *training_end_data*.

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from pandas_datareader import data
```

```
tickers = 'AAPL'
```

```
start_date = '1980-12-01'
end_date = '2018-12-31'
```

```
stock_data = data.get_data_yahoo(tickers, start_date, end_date)
```

```
stock_data.head(10)
```

| | High | Low | Open | Close | Volume | Adj Close |
|------------|----------|----------|----------|----------|-------------|-----------|
| Date | | | | | | |
| 1980-12-12 | 0.515625 | 0.513393 | 0.513393 | 0.513393 | 117258400.0 | 0.023007 |
| 1980-12-15 | 0.488839 | 0.486607 | 0.488839 | 0.486607 | 43971200.0 | 0.021807 |
| 1980-12-16 | 0.453125 | 0.450893 | 0.453125 | 0.450893 | 26432000.0 | 0.020206 |
| 1980-12-17 | 0.464286 | 0.462054 | 0.462054 | 0.462054 | 21610400.0 | 0.020706 |
| 1980-12-18 | 0.477679 | 0.475446 | 0.475446 | 0.475446 | 18362400.0 | 0.021307 |
| 1980-12-19 | 0.506696 | 0.504464 | 0.504464 | 0.504464 | 12157600.0 | 0.022607 |
| 1980-12-22 | 0.531250 | 0.529018 | 0.529018 | 0.529018 | 9340800.0 | 0.023707 |
| 1980-12-23 | 0.553571 | 0.551339 | 0.551339 | 0.551339 | 11737600.0 | 0.024708 |
| 1980-12-24 | 0.582589 | 0.580357 | 0.580357 | 0.580357 | 12000800.0 | 0.026008 |
| 1980-12-26 | 0.636161 | 0.633929 | 0.633929 | 0.633929 | 13893600.0 | 0.028409 |

```
stock_data.describe()
```

| | High | Low | Open | Close | Volume | Adj Close |
|-------|-------------|-------------|-------------|-------------|--------------|-------------|
| count | 9594.000000 | 9594.000000 | 9594.000000 | 9594.000000 | 9.594000e+03 | 9594.000000 |
| mean | 26.549240 | 26.026566 | 26.297032 | 26.292735 | 8.758682e+07 | 22.587725 |
| std | 47.280499 | 46.462657 | 46.880676 | 46.878276 | 8.676287e+07 | 44.664584 |
| min | 0.198661 | 0.196429 | 0.198661 | 0.196429 | 3.472000e+05 | 0.008803 |
| 25% | 1.071429 | 1.031295 | 1.049464 | 1.051429 | 3.407435e+07 | 0.148196 |
| 50% | 1.696429 | 1.633929 | 1.664286 | 1.665179 | 5.944890e+07 | 0.879176 |
| 75% | 26.317500 | 25.404285 | 25.910000 | 25.841429 | 1.091926e+08 | 17.231145 |
| max | 233.470001 | 229.779999 | 230.779999 | 232.070007 | 1.855410e+09 | 230.275482 |

```
stock_data_len = stock_data['Close'].count()
print(stock_data_len)
```

```
9594
```

I'm only interested in *close* prices

```
close_prices = stock_data.iloc[:, 1:2].values
print(close_prices)
```

```
[[ 0.51339287]
 [ 0.48660713]]
```

```
[ 0.45089287]
...
[150.07000732]
[154.55000305]
[156.47999573]]
```

Of course, some of the weekdays might be public holidays in which case no price will be available. For this reason, we will fill the missing prices with the latest available prices

```
all_bussinesdays = pd.date_range(start=start_date, end=end_date, freq='B')
print(all_bussinesdays)
```

```
DatetimeIndex(['1980-12-01', '1980-12-02', '1980-12-03', '1980-12-04',
               '1980-12-05', '1980-12-08', '1980-12-09', '1980-12-10',
               '1980-12-11', '1980-12-12',
               ...
               '2018-12-18', '2018-12-19', '2018-12-20', '2018-12-21',
               '2018-12-24', '2018-12-25', '2018-12-26', '2018-12-27',
               '2018-12-28', '2018-12-31'],
              dtype='datetime64[ns]', length=9936, freq='B')
```

```
close_prices = stock_data.reindex(all_bussinesdays)
close_prices = stock_data.fillna(method='ffill')
```

```
close_prices.head(10)
```

| | High | Low | Open | Close | Volume | Adj Close |
|-------------------|----------|----------|----------|----------|-------------|-----------|
| Date | | | | | | |
| 1980-12-12 | 0.515625 | 0.513393 | 0.513393 | 0.513393 | 117258400.0 | 0.023007 |
| 1980-12-15 | 0.488839 | 0.486607 | 0.488839 | 0.486607 | 43971200.0 | 0.021807 |
| 1980-12-16 | 0.453125 | 0.450893 | 0.453125 | 0.450893 | 26432000.0 | 0.020206 |
| 1980-12-17 | 0.464286 | 0.462054 | 0.462054 | 0.462054 | 21610400.0 | 0.020706 |
| 1980-12-18 | 0.477679 | 0.475446 | 0.475446 | 0.475446 | 18362400.0 | 0.021307 |
| 1980-12-19 | 0.506696 | 0.504464 | 0.504464 | 0.504464 | 12157600.0 | 0.022607 |
| 1980-12-22 | 0.531250 | 0.529018 | 0.529018 | 0.529018 | 9340800.0 | 0.023707 |
| 1980-12-23 | 0.553571 | 0.551339 | 0.551339 | 0.551339 | 11737600.0 | 0.024708 |
| 1980-12-24 | 0.582589 | 0.580357 | 0.580357 | 0.580357 | 12000800.0 | 0.026008 |
| 1980-12-26 | 0.636161 | 0.633929 | 0.633929 | 0.633929 | 13893600.0 | 0.028409 |

The dataset is now complete and free of missing values. Let's have a look to the data frame summary:

Feature scaling

```
training_set = close_prices.iloc[:, 1:2].values
```

```
print(training_set)
```

```
[[ 0.51339287]
 [ 0.48660713]
 [ 0.45089287]
 ...
 [150.07000732]
 [154.55000305]
 [156.47999573]]
```

```
from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler(feature_range = (0, 1))
training_set_scaled = sc.fit_transform(training_set)
print(training_set_scaled.shape)
```

```
(9594, 1)
```

LSTMs expect the data in a specific format, usually a 3D tensor. I start by creating data with 60 days and converting it into an array using NumPy. Next, I convert the data into a 3D dimension array with feature_set samples, 60 days and one feature at each step.

```
features = []
labels = []
for i in range(60, stock_data_len):
    features.append(training_set_scaled[i-60:i, 0])
    labels.append(training_set_scaled[i, 0])

features = np.array(features)
labels = np.array(labels)

features = np.reshape(features, (features.shape[0], features.shape[1], 1))

print(labels)

[0.00082642 0.00089448 0.00087503 ... 0.6528062  0.67231978 0.68072627]
```

```
print(features)

[[[1.38060532e-03]
 [1.26393438e-03]
 [1.10837330e-03]
 ...
 [1.13754103e-03]
 [9.81979819e-04]
 [8.94476613e-04]]

 [1.26393438e-03]
 [1.10837330e-03]
 [1.15698610e-03]
 ...
 [9.81979819e-04]
 [8.94476613e-04]
 [8.26418607e-04]]

 [1.10837330e-03]
 [1.15698610e-03]
 [1.21532157e-03]
 ...
 [8.94476613e-04]
 [8.26418607e-04]
 [8.94476613e-04]]

 ...

 [9.85059938e-01]
 [9.86279533e-01]
 [1.00000000e+00]
 ...
 [6.50889679e-01]
 [6.37648276e-01]
 [6.38214540e-01]]

 [9.86279533e-01]
 [1.00000000e+00]
 [9.86715064e-01]
 ...
 [6.37648276e-01]
 [6.38214540e-01]
 [6.52806203e-01]]

 [1.00000000e+00]
 [9.86715064e-01]
 [9.59927459e-01]
 ...
 [6.38214540e-01]
 [6.52806203e-01]
 [6.72319776e-01]]]
```

Feature tensor with three dimension: features[0] contains the ..., features[1] contains the last 60 days of values and features [2] contains the ...

```
print(features.shape)

(9534, 60, 1)
```

✓ Create the LSTM network

Let's create a sequenced LSTM network with 50 units. Also the net includes some dropout layers with 0.2 which means that 20% of the neurons will be dropped.

```
import tensorflow as tf

model = tf.keras.models.Sequential([
    tf.keras.layers.LSTM(units = 50, return_sequences = True, input_shape = (features.shape[1], 1)),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.LSTM(units = 50, return_sequences = True),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.LSTM(units = 50, return_sequences = True),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.LSTM(units = 50),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(units = 1)
])

WARNING: Logging before flag parsing goes to stderr.
W0414 15:18:15.979501 139980101556096 tf_logging.py:161] <tensorflow.python.keras.layers.recurrent.UnifiedLSTM object at 0x7f4f34285860>
W0414 15:18:16.001110 139980101556096 tf_logging.py:161] <tensorflow.python.keras.layers.recurrent.UnifiedLSTM object at 0x7f4f34285ef0>
W0414 15:18:16.006853 139980101556096 tf_logging.py:161] <tensorflow.python.keras.layers.recurrent.UnifiedLSTM object at 0x7f4f342a6748>
W0414 15:18:16.013553 139980101556096 tf_logging.py:161] <tensorflow.python.keras.layers.recurrent.UnifiedLSTM object at 0x7f4f342a6e48>
```

```
print(model.summary())
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|------------------------------|----------------|---------|
| ===== | | |
| unified_lstm (UnifiedLSTM) | (None, 60, 50) | 10400 |
| dropout (Dropout) | (None, 60, 50) | 0 |
| unified_lstm_1 (UnifiedLSTM) | (None, 60, 50) | 20200 |
| dropout_1 (Dropout) | (None, 60, 50) | 0 |
| unified_lstm_2 (UnifiedLSTM) | (None, 60, 50) | 20200 |
| dropout_2 (Dropout) | (None, 60, 50) | 0 |
| unified_lstm_3 (UnifiedLSTM) | (None, 50) | 20200 |
| dropout_3 (Dropout) | (None, 50) | 0 |
| dense (Dense) | (None, 1) | 51 |
| ===== | | |
| Total params: 71,051 | | |
| Trainable params: 71,051 | | |
| Non-trainable params: 0 | | |
| ===== | | |
| None | | |

The model will be compiled and optimize by the adam optimizer and set the loss function as mean_squared_error

```
model.compile(optimizer = 'adam', loss = 'mean_squared_error')

from time import time
start = time()
history = model.fit(features, labels, epochs = 20, batch_size = 32, verbose = 1)
end = time()

Epoch 1/20
9534/9534 [=====] - 19s 2ms/sample - loss: 0.0017
Epoch 2/20
9534/9534 [=====] - 17s 2ms/sample - loss: 8.3861e-04
Epoch 3/20
9534/9534 [=====] - 17s 2ms/sample - loss: 8.4889e-04
Epoch 4/20
9534/9534 [=====] - 17s 2ms/sample - loss: 6.7823e-04
Epoch 5/20
9534/9534 [=====] - 17s 2ms/sample - loss: 6.0230e-04
```

```

Epoch 6/20
9534/9534 [=====] - 17s 2ms/sample - loss: 5.9530e-04
Epoch 7/20
9534/9534 [=====] - 17s 2ms/sample - loss: 5.6053e-04
Epoch 8/20
9534/9534 [=====] - 18s 2ms/sample - loss: 5.6923e-04
Epoch 9/20
9534/9534 [=====] - 17s 2ms/sample - loss: 4.9843e-04
Epoch 10/20
9534/9534 [=====] - 17s 2ms/sample - loss: 5.7040e-04
Epoch 11/20
9534/9534 [=====] - 17s 2ms/sample - loss: 4.5142e-04
Epoch 12/20
9534/9534 [=====] - 17s 2ms/sample - loss: 5.1318e-04
Epoch 13/20
9534/9534 [=====] - 17s 2ms/sample - loss: 4.6513e-04
Epoch 14/20
9534/9534 [=====] - 17s 2ms/sample - loss: 4.5369e-04
Epoch 15/20
9534/9534 [=====] - 17s 2ms/sample - loss: 4.3363e-04
Epoch 16/20
9534/9534 [=====] - 17s 2ms/sample - loss: 4.8583e-04
Epoch 17/20
9534/9534 [=====] - 17s 2ms/sample - loss: 5.3111e-04
Epoch 18/20
9534/9534 [=====] - 17s 2ms/sample - loss: 4.6322e-04
Epoch 19/20
9534/9534 [=====] - 17s 2ms/sample - loss: 4.5440e-04
Epoch 20/20
9534/9534 [=====] - 17s 2ms/sample - loss: 4.6059e-04

```

```
print('Total training time {} seconds'.format(end - start))
```

```
Total training time 349.84911346435547 seconds
```

```
# [samples, days, features]
print(features.shape)
```

```
(9534, 60, 1)
```

```
testing_start_date = '2019-01-01'
testing_end_date = '2019-04-10'
```

```
test_stock_data = data.get_data_yahoo(tickers, testing_start_date, testing_end_date)
```

```
test_stock_data.tail()
```

| | High | Low | Open | Close | Volume | Adj Close |
|-------------------|------------|------------|------------|------------|------------|------------|
| Date | | | | | | |
| 2019-04-04 | 196.369995 | 193.139999 | 194.789993 | 195.690002 | 19114300.0 | 195.690002 |
| 2019-04-05 | 197.100006 | 195.929993 | 196.449997 | 197.000000 | 18526600.0 | 197.000000 |
| 2019-04-08 | 200.229996 | 196.339996 | 196.419998 | 200.100006 | 25881700.0 | 200.100006 |
| 2019-04-09 | 202.850006 | 199.229996 | 200.320007 | 199.500000 | 35768200.0 | 199.500000 |
| 2019-04-10 | 200.740005 | 198.179993 | 198.679993 | 200.619995 | 21695300.0 | 200.619995 |

```
test_stock_data_processed = test_stock_data.iloc[:, 1:2].values
```

```
print(test_stock_data_processed.shape)
```

```
(69, 1)
```

```
all_stock_data = pd.concat((stock_data['Close'], test_stock_data['Close']), axis = 0)
```

```
inputs = all_stock_data[len(all_stock_data) - len(test_stock_data) - 60:].values
inputs = inputs.reshape(-1,1)
inputs = sc.transform(inputs)
```

```
X_test = []
for i in range(60, 129):
```


```

for i in range(60, 291):
    X_test.append(inputs[i-60:i, 0])

X_test = np.array(X_test)
X_test = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))
predicted_stock_price = model.predict(X_test)
predicted_stock_price = sc.inverse_transform(predicted_stock_price)

plt.figure(figsize=(10,6))
plt.plot(test_stock_data_processed, color='blue', label='Actual Apple Stock Price')
plt.plot(predicted_stock_price , color='red', label='Predicted Apple Stock Price')
plt.title('Apple Stock Price Prediction')
plt.xlabel('Date')
plt.ylabel('Apple Stock Price')
plt.legend()
plt.show()

```



```

test_inputs = test_stock_data_processed.reshape(-1,1)
test_inputs = sc.transform(test_inputs)

print(test_inputs.shape)

test_features = []
for i in range(60, 291):
    test_features.append(test_inputs[i-60:i, 0])

test_features = np.array(test_features)

test_features = np.reshape(test_features, (test_features.shape[0], test_features.shape[1], 1))
print(test_features.shape)

predicted_stock_price = model.predict(test_features)

predicted_stock_price = sc.inverse_transform(predicted_stock_price)
print(predicted_stock_price.shape)

print(test_stock_data_processed.shape)

plt.figure(figsize=(10,6))
plt.plot(test_stock_data_processed, color='blue', label='Actual Apple Stock Price')
plt.plot(predicted_stock_price , color='red', label='Predicted Apple Stock Price')
plt.title('Apple Stock Price Prediction')
plt.xlabel('Date')
plt.ylabel('Apple Stock Price')
plt.legend()
plt.show()

```

